JSPM's
RAJARSHI SHAHU COLLEGE OF ENGINEERING
TATHAWADE, PUNE-33
(An Empowered Autonomous Institute Affiliated to Savitribai Phule Pune University, Pune)

Mini Project

on

"Titanic Survival Prediction using Machine Learning"

-Submitted by-

Aniket B. Pathak **Roll No: CS3223**

Class and Division: TY Div B

Course: Machine Learning

Academic Year 2026-27

(Semester-I)

1. Title

Titanic Survival Prediction using Machine Learning

2. Abstract

The sinking of the RMS Titanic in 1912 caused the death of most of its passengers. Survival was not completely random: passengers of some groups, such as women and first-class travellers, had a much better chance of surviving. This mini project uses **supervised machine learning** to predict whether a Titanic passenger survived or not from details such as ticket class, gender, age, family members on board, fare and port of embarkation.

The dataset of 891 passengers was cleaned, explored with graphs and split into 80% training and 20% testing data. Two classification algorithms, Logistic Regression and AdaBoost, were trained and compared using accuracy, precision, recall and F1-score. **Both models reached 79.9% accuracy** on unseen test data, and Logistic Regression was selected as the final model because of its slightly higher F1-score. Gender and age were found to be the most important factors for survival.

3. Introduction

Machine Learning (ML) is a branch of Artificial Intelligence in which a computer learns patterns from data instead of being programmed with fixed rules. **Supervised learning** is the type of ML in which the model is trained on data that already contains the correct answer (called the label or target). After training, the model can predict the answer for new, unseen data.

When the answer belongs to a fixed set of categories, for example Survived or Died, the task is called **classification**. This project is a binary classification problem: the output is either 1 (survived) or 0 (died).

The Titanic dataset is one of the most popular beginner datasets in machine learning because it is small, easy to understand and contains a mix of numerical and categorical features, missing values and a clear target column.

4. Problem Statement

To build and compare supervised machine learning classification models that predict whether a passenger on the Titanic survived or not, using passenger details such as ticket class, gender, age, number of siblings/spouses and parents/children on board, fare paid and port of embarkation.

5. Objectives

- To understand and apply supervised learning (classification) on a real-world dataset.
- To clean the data by handling missing values and converting text columns into numbers.
- To explore the data using graphs and find which factors affected survival.
- To train two classification models: Logistic Regression and AdaBoost.
- To compare the models using accuracy, precision, recall and F1-score and select the best one.
- To use the best model to predict survival for new passengers.

6. Data set and Libraries required

Data set

The Titanic passenger dataset (file: titanic_raw.csv) contains **891 passengers and 12 columns**. Source: https://github.com/datasciencedojo/datasets (titanic.csv).

Column	Description
Survived	Target: 0 = Died, 1 = Survived
Pclass	Ticket class: 1 (upper), 2 (middle), 3 (lower)
Name, Ticket, PassengerId	Identification details (not useful for prediction)
Sex	Gender of the passenger
Age	Age in years
SibSp	Number of siblings / spouses aboard
Parch	Number of parents / children aboard
Fare	Ticket fare paid
Cabin	Cabin number (mostly missing)
Embarked	Port of embarkation: S = Southampton, C = Cherbourg, Q = Queenstown

Libraries and tools

- **Python 3** with **Jupyter Notebook** as the working environment
- **pandas**: loading and cleaning the data
- **matplotlib** and **seaborn**: drawing graphs
- **scikit-learn**: train-test split, ML models and evaluation metrics

7. Approach/ Methodology

The project follows the standard supervised learning workflow:

Load data → Clean data → Explore (graphs) → Split into train/test → Train models → Evaluate and compare → Predict

- **Data cleaning:** missing Age values were filled with the median age, the 2 missing Embarked values with the most common port, and the Cabin column was dropped because about 77% of it was missing. PassengerId, Name and Ticket were removed as they do not help prediction.
- **Encoding:** machine learning models understand only numbers, so Sex (male = 0, female = 1) and Embarked (S = 0, C = 1, Q = 2) were converted to numeric values.
- **Exploratory data analysis:** bar charts and histograms were used to study survival by gender, class and age.
- **Train-test split:** 80% of the data (712 passengers) was used to train the models and 20% (179 passengers) was kept aside to test them on unseen data.
- **Model training:** Logistic Regression (uses a mathematical formula to separate the two classes) and AdaBoost (builds many small weak models one after another; each new model focuses on the passengers that the earlier models predicted wrongly, and all models are combined by a weighted vote).

- **Evaluation:** Accuracy (correct predictions out of all), Precision (when the model says Survived, how often it is right), Recall (out of all who really survived, how many the model found) and F1-score (balance of precision and recall).

8. Implementation

1. Loading and cleaning the data

```
import pandas as pd
df = pd.read_csv("titanic_raw.csv")

df["Age"] = df["Age"].fillna(df["Age"].median())
df["Embarked"] = df["Embarked"].fillna(df["Embarked"].mode()[0])
df = df.drop(columns=["Cabin"])

df["Sex"] = df["Sex"].map({"male": 0, "female": 1})
df["Embarked"] = df["Embarked"].map({"S": 0, "C": 1, "Q": 2})
```

Column	Missing values	Action taken
Age	177	Filled with median age
Embarked	2	Filled with most common port
Cabin	687	Column dropped

2. Exploratory data analysis

Only about 38% of the passengers survived and 62% died. Women had a survival rate of about 74% compared with about 19% for men. Survival also fell with ticket class: about 63% in first class, 47% in second class and 24% in third class.

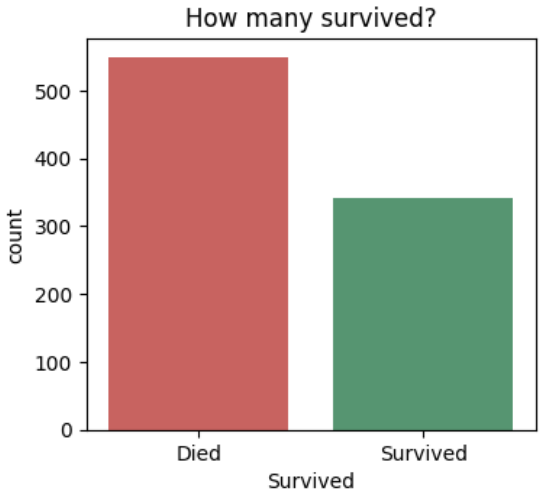


Figure 1: Survival count (0 = Died, 1 = Survived)

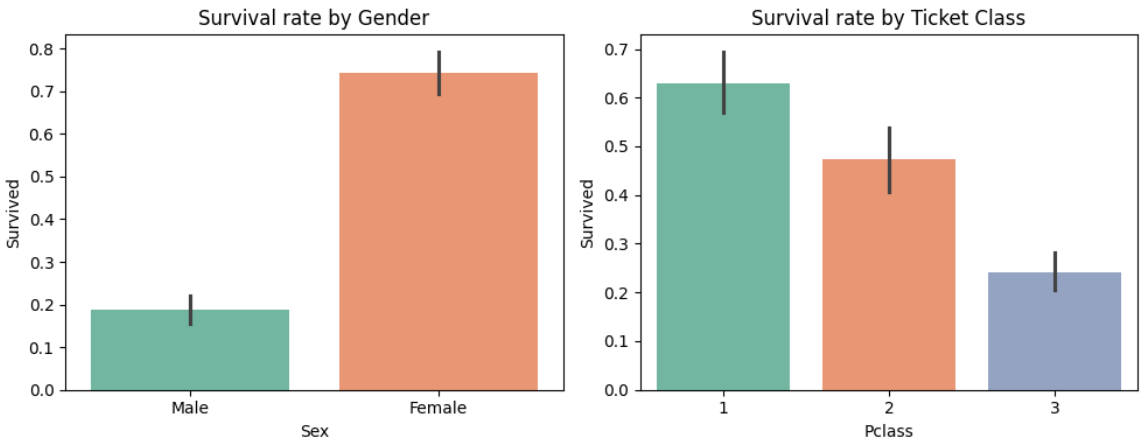


Figure 2: Survival rate by gender and by ticket class

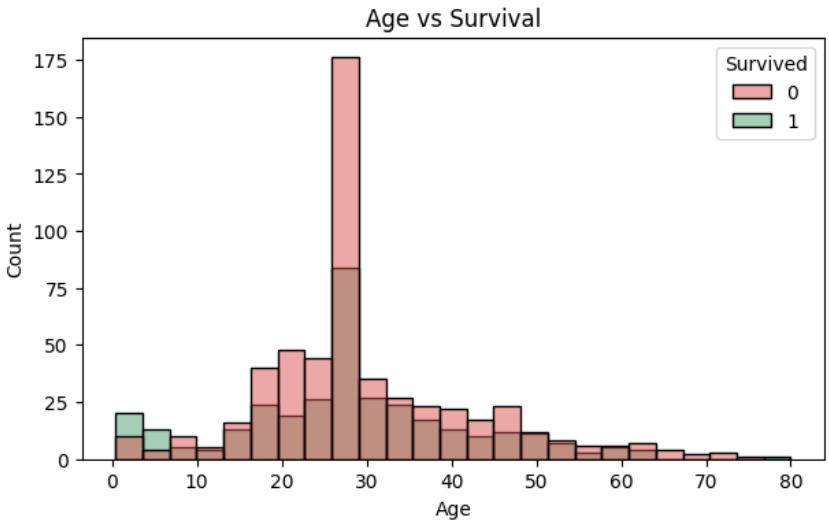


Figure 3: Age distribution of passengers who survived and died

3. Train-test split

```
from sklearn.model_selection import train_test_split

X = df[["Pclass", "Sex", "Age", "SibSp", "Parch", "Fare", "Embarked"]]
y = df["Survived"]
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
```

4. Training and evaluating the models

```
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score, precision_score
from sklearn.metrics import recall_score, f1_score

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "AdaBoost": AdaBoostClassifier(n_estimators=50, random_state=42),
}
for name, model in models.items():
    model.fit(X_train, y_train)          # learn from training data
    pred = model.predict(X_test)         # predict on unseen test data
    print(name, accuracy_score(y_test, pred))
```

9. Results

Both models were tested on the same 179 unseen passengers.

Model	Accuracy	Precision	Recall	F1-Score
Logistic Regression	0.7989	0.7714	0.7297	0.7500
AdaBoost	0.7989	0.7879	0.7027	0.7429

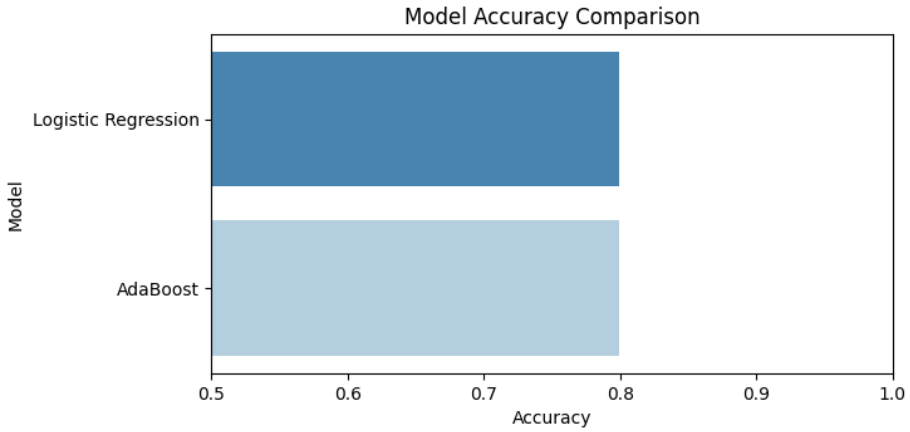


Figure 4: Accuracy comparison of the two models

Both models reached the same accuracy of **79.9%** (143 of the 179 test passengers predicted correctly). Because the accuracy is tied, the F1-score was used to choose the final model: **Logistic Regression (0.75)** is slightly ahead of AdaBoost (0.74). AdaBoost has slightly higher precision, while Logistic Regression has slightly higher recall.

Confusion matrix (Logistic Regression)

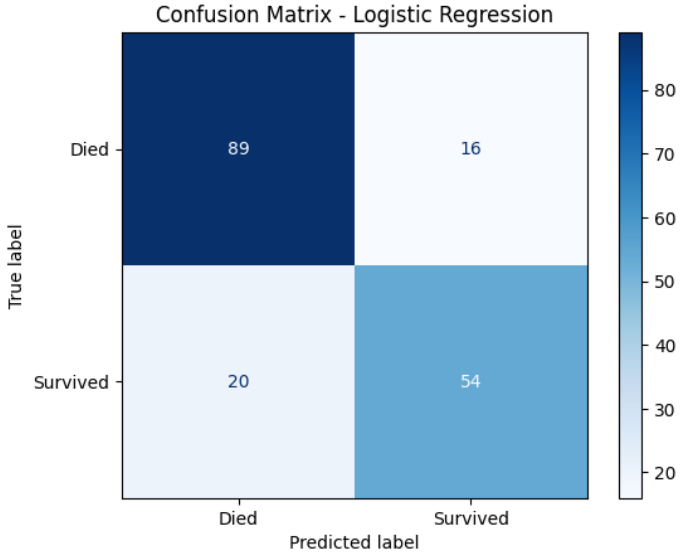


Figure 5: Confusion matrix of the Logistic Regression model

Of the 179 test passengers, 89 who died and 54 who survived were predicted correctly (143 correct). The model wrongly predicted 16 passengers as survived when they died, and 20 as died when they actually survived.

Important features (AdaBoost)

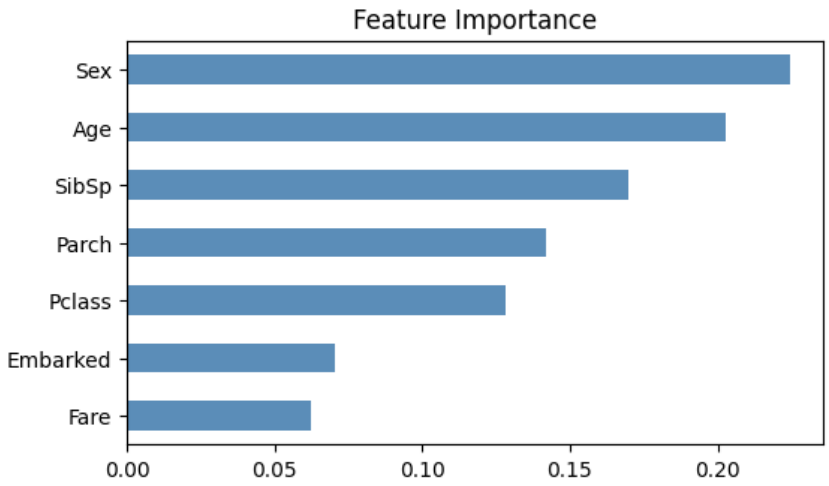


Figure 6: Feature importance from the AdaBoost model

Sex (0.23) and Age (0.20) are the most important features in AdaBoost, followed by SibSp (0.17), Parch (0.14) and Pclass (0.13); Fare (0.06) matters least. In the graphs, gender showed the strongest effect on survival, which matches the "women and children first" rule followed during the evacuation.

Prediction for new passengers

Passenger	Predicted result
1st class woman, age 28	Survived
3rd class man, age 30	Died
2nd class girl, age 8	Survived

10. Conclusion

In this project, supervised machine learning (classification) was used to predict Titanic passenger survival. After cleaning the data and exploring it, two models were trained and tested on unseen passengers. Both **Logistic Regression and AdaBoost reached 79.9% accuracy**; Logistic Regression was selected as the final model because of its slightly higher F1-score (0.75 vs 0.74).

The analysis showed that gender and age are the most important factors, and that women, children and first-class passengers had a much better chance of survival.

Limitations: the dataset is small (891 passengers) and the models were used with basic settings. With only 179 test passengers, one or two passengers change the accuracy slightly. **Future scope:** trying more algorithms such as KNN and SVM, tuning model parameters and creating new features such as family size and passenger title.